

"Meet a FINIX user. It's 11:47 PM, and their phone just buzzed."

SAY: It's 11:47 PM. A FINIX user's phone just buzzed.

SCENE 1 The Phishing SMS

SAY: "Your account will be blocked in 2 hours, click here to verify." This message works on most people because it triggers panic before thought. Watch what FINIX does with it before the user even finishes reading.

DO: Show the SMS, swipe to the fraud badge: Suspicious, in red.

SAY: It's already flagged, not by guessing, but by four checks running the instant it arrived: is the sender actually registered with a bank, is the link known-bad, does the message match this bank's real templates, is it spoofing a bank from a personal number.

FOR THE EXPERTS IN THE ROOM: *A five-layer deterministic rule engine, sender-registry check, URL threat-intel lookup, template-deviation match, spoofing detection, OTP-context validation, with an NLP layer for novel phishing language.*

USP — CONNECTED INTELLIGENCE: *This flag doesn't stay isolated. FINIX runs wealth management and fraud defence off one continuously-learning financial twin of the user. A phishing click, an odd trade, a risky payment, never three alerts, one flagged story. This click just elevated the risk score of the very next transaction, specific to this account.*

SCENE 2 The Near-Miss

SAY: Say the user taps it anyway and it walks them to a transfer.

DO: Pay Anyone, large amount, unfamiliar recipient, Send.

SAY: The moment I hit send, it isn't checking one thing, it's checking fourteen at once, in under a second, unusual amount, first-time recipient, odd hour, impossible device location, and, new even by banking standards, whether this recipient's number is flagged on the government's own fraud registry, the RBI-mandated Financial Fraud Risk Indicator queried live against DoT's Digital Intelligence Platform.

DO: Risk Warning screen: High risk, transaction paused.

FOR THE EXPERTS IN THE ROOM: *A 32-feature ONNX risk model on the transaction, paired with a graph neural network that scores the recipient, how close they sit to a known mule network, a question a rules engine can't ask. Underneath it: device-bound JWT, certificate pinning, rate limiting, idempotency.*

USP — SIM-SWAP-PROOF, ZERO-TRUST: *That JWT rides on FINIX's Unique Binding Token, built to RBI and NPCI's SIM-and-device-binding standard, re-verified on every app open and every transaction. Every session, every API call, every transaction, continuously re-verified, zero implicit trust, anywhere.*

SCENE 3 Why It's Blocked, in Plain Language

SAY: It doesn't just say blocked, it says why, in plain language, right on screen: eight times the usual amount, first-time recipient, unusual hour, elevated registry risk.

SAY: That explanation is generated fresh from the actual signals that fired. We call it explainable AI, every decision can show its work.

FOR THE EXPERTS IN THE ROOM: *One shared explanation contract across every model, SHAP attribution for tree-based models, formula terms for deterministic ones. It cannot hallucinate a reason.*

SCENE 4 Cooling-Off, Not Silence

SAY: It doesn't freeze and walk away. The user gets a cooling-off window and a clear path to confirm, with extra verification, that this really is them.

DO: Show the cooling-off / override screen.

SAY: Under RBI's updated fraud rules, a bank can't unilaterally declare fraud and lock someone out, the account holder gets a real chance to respond. This screen is that chance.

SCENE 5 The Record That Can't Vanish

SAY: Every step, the flag, the score, the explanation, the cooling-off, is already written to an immutable, blockchain-anchored audit log. If this were confirmed fraud, FINIX has a direct path to file with the National Cyber Crime portal and the bank's call centre.

DO: Show the Audit Log / Security Events screen, scrolled to this incident.

FOR THE EXPERTS IN THE ROOM: *Every identity, session, and transaction runs through a hybrid cryptographic architecture, classical encryption plus post-quantum cryptography, hashed to the ledger only after that protection is*

applied. Proof, not a promise. Every approved transaction also gets an independent veto from a smart contract running as chain code on Hyperledger Fabric, outside our application layer, so if our code has a bug, the chain still says no. Then it settles to a double-entry ledger, debits equal credits, every time.

Protected. Now Let's Grow.

SAY: That's FINIX protecting money that's already there. It's a wealth management app that's genuinely serious about security. Now the other 99.9 percent: helping this person build something.

DO: Navigate to the Financial Health Score screen. Let this land.

SCENE 6 The Honest Score

SAY: One number, liquidity, debt, savings, investments, insurance, goals, habits, into a single honest picture.

DO: Scroll to a shock-month badge, e.g. Diwali, marked Excluded.

SAY: This was a big-spend month, Diwali. Most apps would just ding the score. This one recognised the pattern, excluded it, and told the user why. Real financial health isn't about never having an unusual month, it's the pattern underneath.

FOR THE EXPERTS IN THE ROOM: *Median and median-absolute-deviation throughout, not mean and standard deviation, so one outlier month can't distort the score, plus a shock-month detector and a single-step score-change cap.*

SCENE 7 A Goal, Made Real

SAY: I'm creating a real goal....

DO: Goal creation: category, name, target amount, suggested monthly contribution.

SCENE 8 Learning to Invest, Without the Jargon

SAY: Foreclose a loan, or keep investing? Watch this.

DO: What-If Simulation, foreclose vs. continue SIP.

SAY: Real amortisation, real opportunity cost, not a rule of thumb, telling the user which choice wins, and by how much. It will never tell you to buy a specific stock. Numbers and reasoning, you decide.

FOR THE EXPERTS IN THE ROOM: *Every investment-facing model is a transparent, deterministic formula, never a trained recommendation model, staying unambiguously on the informational side of SEBI's investment-advice boundary, auditable by a human planner.*

SCENE 9 Ask It Anything

SAY: For everything else, one place to ask.

DO: AI Chatbot: "How much can I save on tax this year?"

SAY: It answers from this user's real account data, live, held to the same explainability and no-guaranteed-returns rules as everything else. Supports 4 languages.

SCENE 10 The Full Picture

SAY: Tax planning with exact deduction savings. Insurance checked against real dependents. Loans built wisely. One app, one login, one honest picture.

DO: Fast swipe: Tax, Insurance, Loans.

SCENE TRUST The Rules We Never Break

SAY: Every model follows the same Ethical AI Charter: "guaranteed" is a word this AI is programmatically banned from using. No dark patterns. A human can always override, a single two-second press freeze. Data used only for what you consented to, nothing more. Bias-audited every quarter, across income, gender, age, geography.

SAY: Underneath all of it: full RBI, SEBI, IRDAI, and DPDP Act compliance, every formula auditable line by line. This AI advises. The human decides.

SCENE CLOSE Back to 11:47 PM

SAY: Back to 11:47 PM. A phone buzzes with a message built to make someone panic. In most apps, that's where it goes wrong. In FINIX, that's the moment it was built for.

That's FINIX, a trust anyone in the room can verify.